

notebook-5

October 7, 2024

1 Assignment 5: Clustering

1.1 Problem 1: Preprocessing the dataset

Read the dataset:

```
[45]: import itertools

import numpy as np
import pandas as pd
from sklearn.preprocessing import StandardScaler

column_labels = [
    'Area',
    'Perimeter',
    'Compactness',
    'Length of kernel',
    'Width of kernel',
    'Asymmetry coefficient',
    'Length of the kernel groove',
    'Numerical class label',
]
df = pd.read_csv(
    'data/seeds.tsv',
    sep='\t',
    header=None,
    names=column_labels
)
df.head()
```

```
[45]:
```

	Area	Perimeter	Compactness	Length of kernel	Width of kernel	\
0	15.26	14.84	0.8710	5.763	3.312	
1	14.88	14.57	0.8811	5.554	3.333	
2	14.29	14.09	0.9050	5.291	3.337	
3	13.84	13.94	0.8955	5.324	3.379	
4	16.14	14.99	0.9034	5.658	3.562	

Asymmetry coefficient Length of the kernel groove Numerical class label

0	2.221	5.220	1
1	1.018	4.956	1
2	2.699	4.825	1
3	2.259	4.805	1
4	1.355	5.175	1

Scale the dataset using a normalizer:

```
[46]: scaler = StandardScaler()

scaled_data = pd.DataFrame(
    scaler.fit_transform(
        df.drop(columns="Numerical class label")
    ),
    columns=column_labels[:-1]
)
scaled_df = pd.concat(
    [scaled_data, df["Numerical class label"]],
    axis=1
)
scaled_df.head()
```

```
[46]:      Area  Perimeter  Compactness  Length of kernel  Width of kernel  \
0  0.142098  0.215462    0.000061      0.304218      0.141702
1  0.011188  0.008224    0.428515     -0.168625      0.197432
2 -0.192067 -0.360201    1.442383     -0.763637      0.208048
3 -0.347091 -0.475333    1.039381     -0.688978      0.319508
4  0.445257  0.330595    1.374509      0.066666      0.805159

      Asymmetry coefficient  Length of the kernel groove  Numerical class label
0                -0.986152                -0.383577                1
1                -1.788166                -0.922013                1
2                -0.667479                -1.189192                1
3                -0.960818                -1.229983                1
4                -1.563495                -0.475356                1
```

1.2 Problem 2: Determining the appropriate number of clusters

Iterating over a range of cluster numbers, determine the optimal number of clusters using the elbow method:

```
[47]: from sklearn.cluster import KMeans
import matplotlib.pyplot as plt

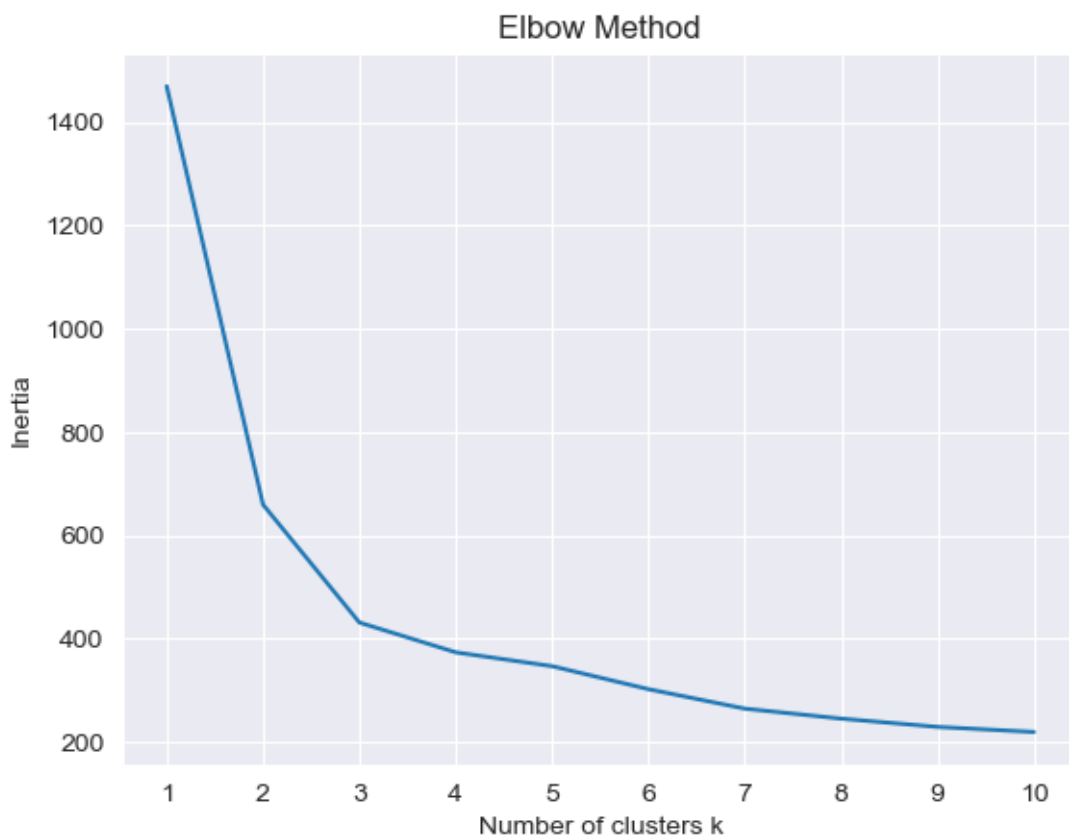
k = range(1, 11)
inertia = []
for i in k:
    kmeans = KMeans(
```

```

        n_clusters=i,
        random_state=0,
        n_init="auto"
    ).fit(scaled_data)
    inertia.append(kmeans.inertia_)

plt.plot(k, inertia)
plt.title('Elbow Method')
plt.xlabel('Number of clusters k')
plt.ylabel('Inertia')
plt.xticks(k)
plt.savefig("figures/elbow_method.png")

```



We can see that the optimal number of clusters is 3.


1.3 Problem 3: Visualizing the classes

Scatterplot the pairs of features, coloring the points according to the cluster they belong to:

```
[76]: import seaborn as sns

sns.pairplot(
    scaled_df,
    hue="Numerical class label",
    height=4
)

plt.title("Pairplot of the features")
plt.savefig("figures/features_pairplot.png")
```




Scatterplot of one particular pair of features:

```
[49]: sns.scatterplot(
    x="Area",
    y="Perimeter",
```

```

    data=scaled_df,
    hue="Numerical class label",
)
plt.title("Area vs Perimeter")
plt.savefig("figures/area_perimeter_scatterplot.png")

```



Scatterplot the Gaussian random projections of the dataset, coloring the points according to the cluster they belong to:

```

[50]: from sklearn.random_projection import \
        GaussianRandomProjection

embedding = GaussianRandomProjection(
    random_state=42, n_components=2
).fit_transform(scaled_data)


sns.scatterplot(
    x=embedding[:, 0],
    y=embedding[:, 1],
    hue=df["Numerical class label"]
)

```

```

)
plt.title("Gaussian Random Projection of the dataset")
plt.xlabel("Component 1")
plt.ylabel("Component 2")
plt.savefig(
    "figures/gaussian_random_projection.png"
)

```



Scatterplot the UMAP projections of the dataset, coloring the points according to the cluster they belong to:

```

[51]: from umap import UMAP

embedding = UMAP(
    random_state=42, n_components=2
).fit_transform(scaled_data)

sns.scatterplot(
    x=embedding[:, 0],
    y=embedding[:, 1],

```


```

    hue=df["Numerical class label"]
)
plt.title("UMAP projection of the dataset")
plt.xlabel("Component 1")
plt.ylabel("Component 2")
plt.savefig("figures/umap_projection.png")

```

/Users/giacomo/.miniforge/envs/ds/lib/python3.11/site-packages/umap/umap_.py:1945: UserWarning: n_jobs value 1 overridden to 1 by setting random_state. Use no seed for parallelism.

warn(f"n_jobs value {self.n_jobs} overridden to 1 by setting random_state. Use no seed for parallelism.")



1.4 Problem 4: Evaluating clustering

Cluster the dataset using the optimal number of clusters:

```

[52]: k = scaled_df["Numerical class label"].nunique()

kmeans = KMeans(
    n_clusters=k,

```

```

        random_state=0,
        n_init="auto"
    ).fit(scaled_data)

```

Compute the rand index:

```

[53]: from sklearn.metrics import rand_score, \
        accuracy_score

rand_index = rand_score(
    scaled_df["Numerical class label"],
    kmeans.labels_
)
print(f"Rand index: {rand_index:.3f}")

```

Rand index: 0.900

Compute the accuracy by finding the maximum accuracy on all the cluster labels permutations:

```

[54]: def compute_accuracy(y_true, y_pred):
        max_accuracy = float('-inf')
        combinations = list(
            itertools.permutations(range(1, k + 1))
        )
        for combination in combinations:
            transformed_labels = np.array(
                [combination[label] for label in y_pred]
            )
            max_accuracy = max(
                max_accuracy,
                accuracy_score(y_true, transformed_labels)
            )
        return max_accuracy

accuracy = compute_accuracy(
    scaled_df["Numerical class label"],
    kmeans.labels_
)
print(f"Accuracy: {accuracy:.3f}")

```

Accuracy: 0.919

1.4.1 Problem 5: Agglomerative clustering

Hierarchically cluster the dataset with all the linkage methods:

```

[55]: from sklearn.cluster import \
        AgglomerativeClustering

```



```

best_method = None
best_accuracy = float('-inf')
linkage_methods = [
    "ward",
    "complete",
    "average",
    "single"
]
for method in linkage_methods:
    cluster = AgglomerativeClustering(
        n_clusters=k,
        linkage=method
    ).fit(scaled_data)
    accuracy = compute_accuracy(
        scaled_df["Numerical class label"],
        cluster.labels_
    )

    if accuracy > best_accuracy:
        best_accuracy = accuracy
        best_method = method

    print(
        f"Accuracy for linkage method {method}: {accuracy:.3f}"
    )

print(f"Best linkage method: {best_method}")

```

Accuracy for linkage method ward: 0.929
 Accuracy for linkage method complete: 0.876
 Accuracy for linkage method average: 0.881
 Accuracy for linkage method single: 0.348
 Best linkage method: ward

Plot the dendrogram for the best linkage method:

```

[56]: from scipy.cluster.hierarchy import dendrogram, \
        linkage

plt.figure(figsize=(25, 10))
dendrogram(
    linkage(scaled_data, method=best_method),
    labels=scaled_df[
        "Numerical class label"
    ].values,
    leaf_font_size=10,
    leaf_rotation=0
)

```

```
)
plt.title("Dendrogram of the ward hierarchical clustering")
plt.savefig("figures/dendrogram.png")
```

